

# Druid for real-time analysis

Yann Esposito

7 Avril 2016

# Druid the Sales Pitch

# Intro

# Experience

- ▶ Real Time Social Media Analytics

# Real Time?

- ▶ Ingestion Latency: seconds
- ▶ Query Latency: seconds

# Demand

- ▶ Twitter: 20k msg/s, 1msg = 10ko during 24h
- ▶ Facebook public: 1000 to 2000 msg/s continuously
- ▶ Low Latency

# Reality

- ▶ Twitter: 400 msg/s continuously, burst to 1500
- ▶ Facebook: 1000 to 2000 msg/s

# Origin (PHP)



# 1st Refactoring (Node.js)

# Return of Experience

# Return of Experience

## 2nd Refactoring

## 2nd Refactoring (FTW!)

## 2nd Refactoring return of experience

# Demo

# Pre Considerations



# Discovered vs Invented

Try to conceptualize a s.t.

- ▶ Ingest Events
- ▶ Real-Time Queries
- ▶ Scalable
- ▶ Highly Available

Analytics: timeseries, alerting system, top N, etc...

# In the End

Druid concepts are always emerging naturally

# Druid

# Who?

Metamarkets

Powered by Druid

- ▶ Alibaba, Cisco, Criteo, eBay, Hulu, Netflix, Paypal...

# Goal

*Druid is an open source store designed for real-time exploratory analytics on large data sets.*

*hosted dashboard that would allow users to arbitrarily explore and visualize event streams.*

# Concepts

- ▶ Column-oriented storage layout
- ▶ distributed, shared-nothing architecture
- ▶ advanced indexing structure

# Key Features

- ▶ Sub-second OLAP Queries
- ▶ Real-time Streaming Ingestion
- ▶ Power Analytic Applications
- ▶ Cost Effective
- ▶ High Available
- ▶ Scalable

# Right for me?

- ▶ require fast aggregations
- ▶ exploratory analytics
- ▶ analysis in real-time
- ▶ lots of data (trillions of events, petabytes of data)
- ▶ no single point of failure



# High Level Architecture

# Inspiration

- ▶ Google's BigQuery/Dremel
- ▶ Google's PowerDrill

# Index / Immutability

Druid indexes data to create mostly immutable views.

# Storage

Store data in custom column format highly optimized for aggregation & filter.

# Specialized Nodes

- ▶ A Druid cluster is composed of various type of nodes
- ▶ Each designed to do a small set of things very well
- ▶ Nodes don't need to be deployed on individual hardware
- ▶ Many node types can be colocated in production

# Druid vs X

# Elasticsearch

- ▶ resource requirement much higher for ingestion & aggregation
- ▶ No data summarization (100x in real world data)

# Key/Value Stores (HBase/Cassandra/OpenTSDB)

- ▶ Must Pre-compute Result
  - ▶ Exponential storage
  - ▶ Hours of pre-processing time
- ▶ Use the dimensions as key (like in OpenTSDB)
  - ▶ No filter index other than range
  - ▶ Hard for complex predicates



# Spark

- ▶ Druid can be used to accelerate OLAP queries in Spark
- ▶ Druid focuses on the latencies to ingest and serve queries
- ▶ Too long for end user to arbitrarily explore data

# SQL-on-Hadoop (Impala/Drill/Spark SQL/Presto)

- ▶ Queries: more data transfer between nodes
- ▶ Data Ingestion: bottleneck by backing store
- ▶ Query Flexibility: more flexible (full joins)

# Data

# Concepts

- ▶ **Timestamp column:** query centered on time axis
- ▶ **Dimension columns:** strings (used to filter or to group)
- ▶ **Metric columns:** used for aggregations (count, sum, mean, etc...)

# Indexing

- ▶ Immutable snapshots of data
- ▶ data structure highly optimized for analytic queries
- ▶ Each column is stored separately
- ▶ Indexes data on a per shard (segment) level

# Loading

- ▶ Real-Time
- ▶ Batch

# Querying

- ▶ JSON over HTTP
- ▶ Single Table Operations, no joins.

# Segments

- ▶ Per time interval
  - ▶ skip segments when querying
- ▶ Immutable
  - ▶ Cache friendly
  - ▶ No locking
- ▶ Versioned
  - ▶ No locking
  - ▶ Read-write concurrency



# Roll-up

# Example

timestamp	page	...	added	deleted
2011-01-01T00:01:35Z	Cthulhu		10	65
2011-01-01T00:03:63Z	Cthulhu		15	62
2011-01-01T01:04:51Z	Cthulhu		32	45
2011-01-01T01:01:00Z	Azatoth		17	87
2011-01-01T01:02:00Z	Azatoth		43	99
2011-01-01T02:03:00Z	Azatoth		12	53

timestamp	page	...	nb	added	deleted
2011-01-01T00:00:00Z	Cthulhu		2	25	127
2011-01-01T01:00:00Z	Cthulhu		1	32	45
2011-01-01T01:00:00Z	Azatoth		2	60	186
2011-01-01T02:00:00Z	Azatoth		1	12	53

# as SQL

```
GROUP BY timestamp, page, nb, added, deleted  
:: nb = COUNT(1)  
  , added = SUM(added)  
  , deleted = SUM(deleted)
```

In practice can dramatically reduce the size (up to  
x100)

# Segments

# Sharding

```
sampleData_2011-01-01T01:00:00:00Z_2011-01-01T01:00:00:00Z
```

timestamp	page	...	nb	added	deleted
2011-01-01T01:00:00Z	Cthulhu		1	20	45
2011-01-01T01:00:00Z	Azatoth		1	30	106

```
sampleData_2011-01-01T01:00:00:00Z_2011-01-01T01:00:00:00Z
```

timestamp	page	...	nb	added	deleted
2011-01-01T01:00:00Z	Cthulhu		1	12	45
2011-01-01T01:00:00Z	Azatoth		2	30	80

# Core Data Structure

Timestamp	Dimensions				Metrics	
Timestamp	Page	Username	Gender	City	Characters Added	Characters Removed
2011-01-01T01:00:00Z	Justin Bieber	Boxer	Male	San Francisco	1800	25
2011-01-01T01:00:00Z	Justin Bieber	Reach	Male	Waterloo	2912	42
2011-01-01T02:00:00Z	Ke\$ha	Helz	Male	Calgary	1953	17
2011-01-01T02:00:00Z	Ke\$ha	Xeno	Male	Taiyuan	3194	170

- ▶ dictionary
- ▶ a bitmap for each value
- ▶ a list of the columns values encoded using the dictionary

# Example

```
dictionary: { "Cthulhu": 0  
              , "Azatoth": 1 }
```

```
column data: [0, 0, 1, 1]
```

```
bitmaps (one for each value of the column):  
value="Cthulhu": [1,1,0,0]  
value="Azatoth": [0,0,1,1]
```

# Example (multiple matches)

```
dictionary: { "Cthulhu": 0  
              , "Azatoth": 1 }
```

```
column data: [0, [0,1], 1, 1]
```

bitmaps (one for each value of the column):

```
value="Cthulhu": [1,1,0,0]
```

```
value="Azatoth": [0,1,1,1]
```



# Real-time ingestion

- ▶ Via Real-Time Node and Firehose
  - ▶ No redundancy or HA, thus not recommended
- ▶ Via Indexing Service and Tranquility API
  - ▶ Core API
  - ▶ Integration with Streaming Frameworks
  - ▶ HTTP Server
  - ▶ **Kafka Consumer**

# Batch Ingestion

- ▶ File based (HDFS, S3, ...)

# Real-time Ingestion

Task 1: [ Interval ] [ Window ]

Task 2: [ ]

time

# Querying

# Query types

- ▶ Group by: group by multiple dimensions
- ▶ Top N: like grouping by a single dimension
- ▶ Timeseries: without grouping over dimensions
- ▶ Search: Dimensions lookup
- ▶ Time Boundary: Find available data timeframe
- ▶ Metadata queries

# Example(s)

```
{ "queryType": "groupBy",  
  "dataSource": "druidtest",  
  "granularity": "all",  
  "dimensions": [],  
  "aggregations": [  
    { "type": "count", "name": "rows" },  
    { "type": "longSum", "name": "imps", "fieldName": "imp" },  
    { "type": "doubleSum", "name": "wp", "fieldName": "wp" }  
  ],  
  "intervals": ["2010-01-01T00:00/2020-01-01T00:00"] }
```

# Result

```
[ {  
  "version" : "v1",  
  "timestamp" : "2010-01-01T00:00:00.000Z",  
  "event" : {  
    "imps" : 5,  
    "wp" : 15000.0,  
    "rows" : 5  
  }  
} ]
```

# Caching

- ▶ Historical node level
  - ▶ By segment
- ▶ Broker Level
  - ▶ By segment and query
  - ▶ groupBy is disabled on purpose!
- ▶ By default: local caching



# Druid Components

# Druid

- ▶ Real-time Nodes
- ▶ Historical Nodes
- ▶ Broker Nodes
- ▶ Coordinator
- ▶ For indexing:
  - ▶ Overlord
  - ▶ Middle Manager

# Also

- ▶ Deep Storage (S3, HDFS, ...)
- ▶ Metadata Storage (SQL)
- ▶ Load Balancer
- ▶ Cache

# Coordinator

- ▶ Real-time Nodes (pull data, index it)
- ▶ Historical Nodes (keep old segments)
- ▶ Broker Nodes (route queries to RT & Hist. nodes, merge)
- ▶ Coordinator (manage segments)
- ▶ For indexing:
  - ▶ Overlord (distribute task to the middle manager)
  - ▶ Middle Manager (execute tasks via Peons)

## When *not* to choose Druid

# Graphite (metrics)

# Pivot (exploring data)

# Caravel



# Conclusions

# Precompute your time series?



# Don't reinvent it

- ▶ need a user facing API
- ▶ need time series on many dimensions
- ▶ need real-time
- ▶ big volume of data

# Druid way is the right way!

1. Push in kafka
2. Add the right dimensions
3. Push in druid
4. ???
5. Profit!